



# *PREDICT WHETHER INCOME EXCEEDS \$50K/YR BASED ON CENSUS DATA*

Binary classification using 5  
different machine learning  
algorithms.

Full codes in:

[https://www.kaggle.com/thalia18/binary-  
classification](https://www.kaggle.com/thalia18/binary-classification)

*The prediction task is to determine whether a person makes over \$50K a year.*

I used five different classification algorithms:

- **Decision Tree Classifier**
- **Random Forest Classifier**
- **Logistic classifier**
- **SVM classifier**
- **K Neighbors Classifier**

I evaluated my predictions using different metrics:

- **Accuracy**
- **Precision**
- **Recall**
- **F1**
- **Area under precision recall**



# THE DATA:

This data was extracted from the 1994 Census bureau database.

Attributes:

- >50K, <=50K
- age: continuous
- work class: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, ...
- education: Bachelors, Some-college, Masters, Doctorate, 5th-6th, Preschool...
- education-num: continuous
- marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse
- occupation: Tech-support, Craft-repair, Machine-op-inspct, ...
- relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, ..
- race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black
- sex: Female, Male
- capital-gain: continuous
- capital-loss: continuous
- hours-per-week: continuous
- native-country: United-States, Cambodia, England, ..



- Sample of numerical data:


	age	fnlwgt	education_num	capital_gain	capital_loss	hours_per_week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000


- After cleaning the data, remove null or repeated rows, and transform categorical attributes to numeric.


	age	sex	education_num	hours_per_week	born_usa	white	black	single	married	separated	divorced	widowed	highdegree	capital_gain
0	90	0	9	40	1	1	0	0	0	0	0	1	0	0
1	82	0	9	18	1	1	0	0	0	0	0	1	0	0
2	66	0	10	40	1	0	1	0	0	0	0	1	0	0
3	54	0	4	40	1	1	0	0	0	0	1	0	0	0
4	41	0	10	40	1	1	0	0	0	1	0	0	0	0




# SOME DATA VISUALIZATIONS:

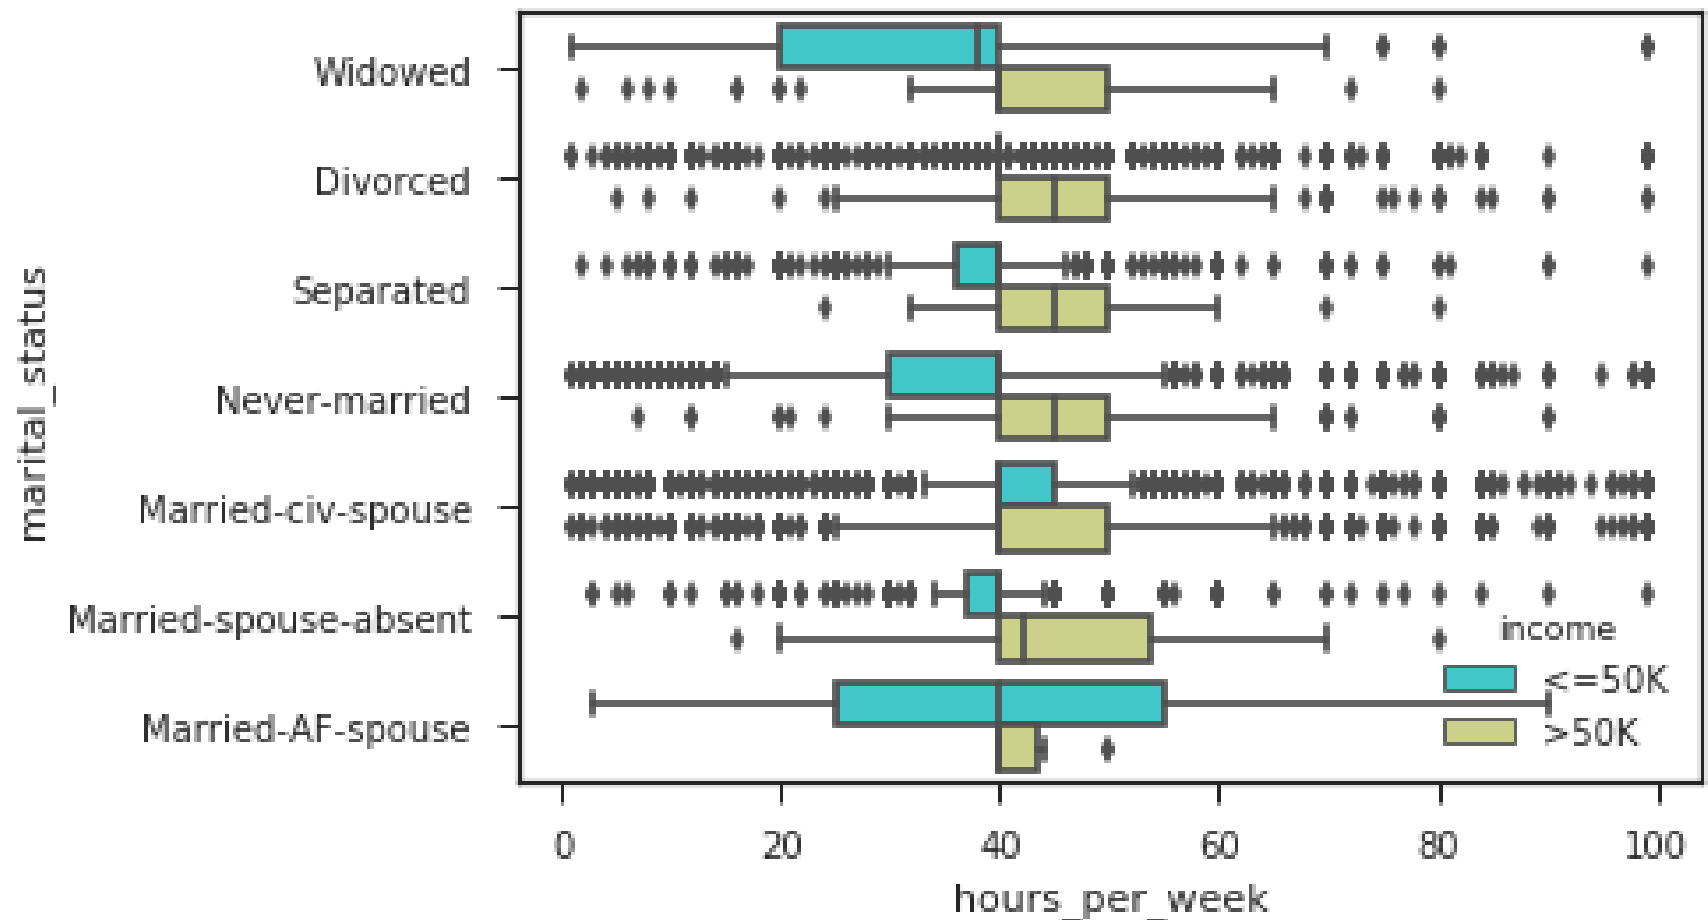














education



I removed attributes that contain just minor categories. I kept attributes that have large categories, for example for race white and black are large categories and for native-country United States is the main category.



So my final attributes are:  
income, age,  
education-num,  
marital-status,  
sex, capital-gain,  
capital-loss, hours  
per week, native  
country.



# MACHINE LEARNING MODELS:

## DECISION TREE:


First I looked for the depth that gives the best accuracy:

Max Depth	Average Accuracy
1	0.759205
2	0.819361
3	0.818102
4	0.818194
5	0.821081
6	0.803427
7	0.826823
8	0.812485
9	0.820805
10	0.813958
11	0.816721
12	0.815216

- The best depth was 7. The full metrics results for this depth is:

```
Accuracy: 0.8529664660361135  
Precision: 0.7872991583779648  
Recall: 0.5282340862422998  
F1: 0.632258064516129  
Area under precision Recall: 0.5287636464397456
```

- We can visualize our decision tree:



# RANDOM FOREST:

```
Accuracy: 0.8349097162510748  
Precision: 0.6928480204342273  
Recall: 0.5569815195071869  
F1: 0.6175298804780875  
Area under precision Recall: 0.49191017630076017
```

# LOGISTIC REGRESSION:

```
Accuracy: 0.8399459525856774  
Precision: 0.7207392197125256  
Recall: 0.5405544147843943  
F1: 0.6177764740393077  
Area under precision Recall: 0.4995361212572641
```



## SVM CLASSIFIER:

Accuracy: 0.8416656430413954

Precision: 0.7716405605935697

Recall: 0.4804928131416838

F1: 0.5922176526415691

Area under precision Recall: 0.49507679663572923

## K NEIGHBORS CLASSIFIER:

Accuracy: 0.8451050239528314

Precision: 0.708308065494239

Recall: 0.5995893223819302

F1: 0.6494300806227412

Area under precision Recall: 0.5205052784173478



# METRICS:

- **Accuracy** Fraction of predictions our model got right

**The best was the classification tree.**

- **Precision** Proportion of those predicted positive, how many of them are actual positive.

**The best was the classification tree.**

- **Recall** Proportion of the actual positive that were predicted correctly

**The best was the Kneighbors.**

- **F1** Conveys the balance between the precision and the recall

**The best was the Kneighbors**

- **Area under precision recall** The precision-recall curve shows the tradeoff between precision and recall for different threshold.

**The best was the classification tree.**

